


serde for SERious DEbugging

Lucas Pardue

HTTP is increasingly binary and multiplexed

HTTP/1x vs. HTTP/2 vs. HTTP/3

Distribution of traffic by HTTP version ?



Data from
<https://radar.cloudflare.com/adoption-and-usage?range=28d>
captured on October 31, 2022

ERR_SPDY_PROTOCOL_ERROR

ERR_HTTP2_PROTOCOL_ERROR

ERR_QUIC_PROTOCOL_ERROR

Question to the room:

What kinds of behaviours in HTTP/2 or HTTP/3 tend to lead to these sorts of problems?

“Logs or it didn’t happen”

Every dev, ever and always

qlog - structured logging by endpoints

Implementations often have logging that can enhance or augment packet captures.

A common logging format can encourage an ecosystem of analysis tools. E.g. what is an endpoint producing and why is it doing that?

[draft-ietf-quic-qlog-main-schema](#): a base schema defined in Concise Data Definition Language (CDDL; [RFC 8610](#)). Highly extensible. Many possible serialization formats.

[Draft-ietf-quic-qlog-quic-events](#), [draft-ietf-quic-qlog-h3-events](#): concrete definitions to cover events related to packets and frames, security, congestion control etc.

qlog CDDL examples

```
HTTPFrameCreated = {  
    stream_id: uint64  
    ? length: uint64  
    frame: $HTTPFrame  
    ? raw: RawInfo  
}
```

```
HTTPFrameParsed = {  
    stream_id: uint64  
    ? length: uint64  
    frame: $HTTPFrame  
    ? raw: RawInfo  
}
```

```
; The HTTPFrame is any key-value map (e.g., JSON object)  
$HTTPFrame /= {  
    * text => any  
}
```

```
$HTTPFrame /= HTTPBaseFrames
```

```
HTTPBaseFrames = HTTPDataFrame / HTTPHeadersFrame /  
    HTTPCancelPushFrame / HTTPSettingsFrame /  
    HTTPPushPromiseFrame / HTTPGoawayFrame /  
    HTTPMaxPushIDFrame / HTTPReservedFrame /  
    HTTPUnknownFrame
```

```
HTTPHeadersFrame = {  
    frame_type: "headers"  
    headers: [* HTTPField]  
}
```

```
HTTPField = {  
    name: text  
    value: text  
}
```

qlog example

Client: QLOGDIR=qlogs quiche-client --no-verify --wire-version 1

<https://127.0.0.1:4433/index.html>

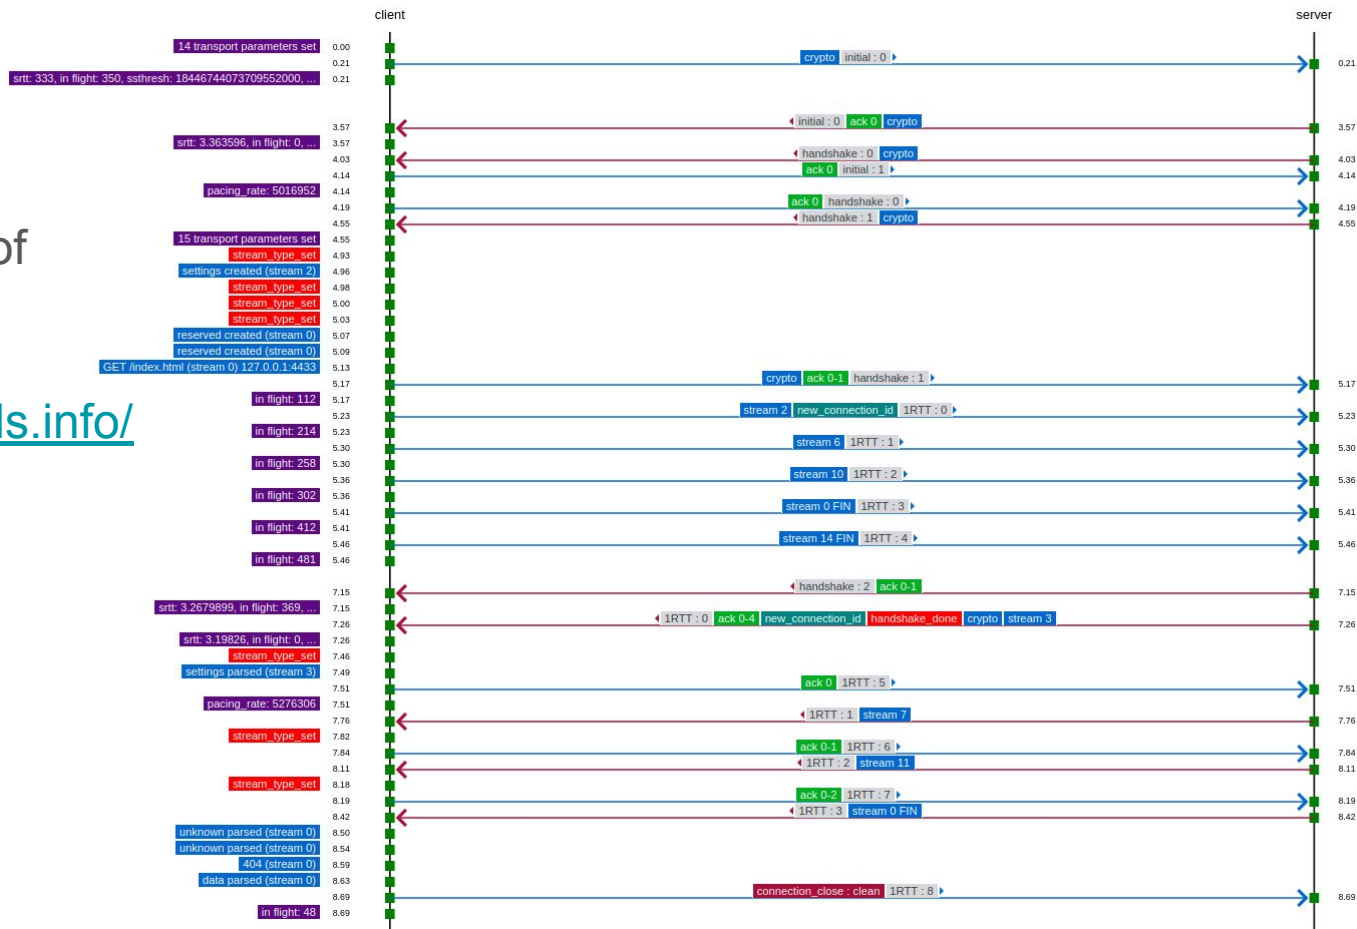
Server: QLOGDIR=qlogs quiche-server --no-retry

```
{"qlog_version":"0.3","qlog_format":"JSON-SEQ","title":"quiche-client qlog","description":"quiche-client qlog  
id=9463b9d6695a7b2d189da2871fc255977bc7c6f8","trace":{"vantage_point":{"type":"client"},"title":"quiche-client  
qlog","description":"quiche-client qlog id=9463b9d6695a7b2d189da2871fc255977bc7c6f8","configuration":{"time_offset":0.0}}}  
{"time":0.0,"name":"transport:parameters_set","data":{"owner":"local","tls_cipher":"None","disable_active_migration":true,"max_idle_t  
imeout":30000,"max_udp_payload_size":1350,"ack_delay_exponent":3,"max_ack_delay":25,"active_connection_id_limit":2,"initial_max_data"  
:10000000,"initial_max_stream_data_bidi_local":1000000,"initial_max_stream_data_bidi_remote":1000000,"initial_max_stream_data_uni":10  
00000,"initial_max_streams_bidi":100,"initial_max_streams_uni":100}}  
{"time":0.207949,"name":"transport:packet_sent","data":{"header":{"packet_type":"initial","packet_number":0,"version":"1","scil":20,"  
dcil":16,"scid":"9463b9d6695a7b2d189da2871fc255977bc7c6f8","dcid":"6c94d2c299cbff6253a202bcb20ceb42"},"raw":{"length":350,"payload_le  
ngth":287},"send_at_time":0.207949,"frames":[{"frame_type":"crypto","offset":0,"length":283}]}}  
{"time":0.207949,"name":"recovery:metrics_updated","data":{"smoothed_rtt":333.0,"rtt_variance":166.5,"congestion_window":13500,"bytes  
_in_flight":350,"ssthresh":18446744073709551615}}  
{"time":3.5715451,"name":"transport:packet_received","data":{"header":{"packet_type":"initial","packet_number":0,"version":"1","scil"  
:20,"dcil":20,"scid":"78015def011d1adf3af94c44067955dd4d52fc70","dcid":"9463b9d6695a7b2d189da2871fc255977bc7c6f8"},"raw":{"length":12  
00,"payload_length":117},"frames":[{"frame_type":"ack","ack_delay":0.305,"acked_ranges":[[0,0]],{"frame_type":"crypto","offset":0,"l  
ength":90}]}}
```

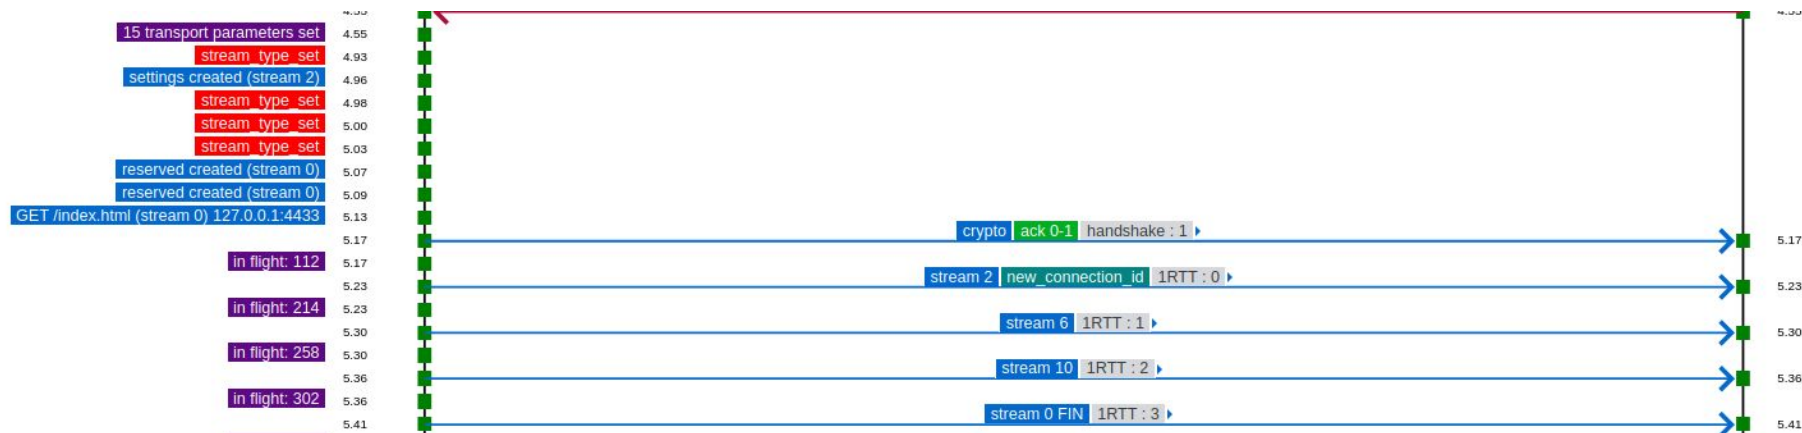

qvis

Making sense out of
oodles of data

<https://qvis.quictools.info/>



Streams example: HTTP/3



Control stream on ID 2. QPACK streams on ID 6 and 10.

Request stream on ID 0. GET request for /index.html. Stream is FIN'd to indicate request message is complete


Streams example: HTTP/3

```
5 0.004963775 127.0.0.1 43959 127.0.0.1 4433 HTTP3 256 Protected Payload (KP0), DCID=78015def011d1adf3af94c44067955dd4d52fc70, PKN: 0, NCI, STREAM(2), SETTINGS
[Packet Length: 102]
  QUIC Short Header DCID=78015def011d1adf3af94c44067955dd4d52fc70 PKN=0
    0... .. = Header Form: Short Header (0)
    .1.. .. = Fixed Bit: True
    ..0. .. = Spin Bit: False
    ...0 0... = Reserved: 0
    ....0... = Key Phase Bit: False
    .... ..00 = Packet Number Length: 1 bytes (0)
    Destination Connection ID: 78015def011d1adf3af94c44067955dd4d52fc70
    Packet Number: 0
    Protected Payload: 03b7d8dfe40be2186a8251313d79001ec5d1d0e10dc73ae1213658fe7cfa6292b991553f...
  NEW_CONNECTION_ID
    Frame Type: NEW_CONNECTION_ID (0x0000000000000018)
    Sequence: 1
    Retire Prior To: 0
    Connection ID Length: 20
    Connection ID: 5a5896ac2c7ba6d164c6b616bd6409af74edd55f
    Stateless Reset Token: 86a804a6b016cc69312dce777734425b
  STREAM id=2 fin=0 off=0 len=19 uni=1
    Frame Type: STREAM (0x000000000000000e)
    Stream ID: 2
    Offset: 0
    Length: 19
    Stream Data: 000410e0b9395476f5e936ef7147d23285d941
  Hypertext Transfer Protocol Version 3
    Stream type: Control Stream (0x0000000000000000)
    Type: SETTINGS (0x0000000000000004)
    Length: 16
    Frame Payload: e0b9395476f5e936ef7147d23285d941
```


Example: HTTP/3 prioritization shown in qvis

5 concurrent transfers of 5 MB, all urgency=1

quiche (before priorities)
round-robin



quiche (now)
FIFO



qvis is great (but has some limitations)

Browser-based tool means it's instantly accessible

Large qlogs slow to load or process. Too large they'll crash the browser process

Qvis is quite an interactive tool. Good for point investigations. Less good for bulk investigations over many files.

Qlog is usually just JSON or JSON-SEQ. Pre-processing with tools like jq can help a lot but are not schema aware.

Quiche qlogs

<https://crates.io/crates/qlog> - A general rust crate for serialization and deserialization of qlog. Quiche uses this for qlog serialization to JSON-SEQ. Firefox uses the crate too.

CDDL

```
HTTPFrameCreated = {  
  stream_id: uint64  
  ? length: uint64  
  frame: $HTTPFrame  
  ? raw: RawInfo  
}
```

Rust

```
#[serde_with::skip_serializing_none]  
#[derive(Serialize, Deserialize, Clone,  
PartialEq, Eq, Debug)]  
pub struct H3FrameCreated {  
  pub stream_id: u64,  
  pub length: Option<u64>,  
  pub frame: Http3Frame,  
  pub raw: Option<RawInfo>,  
}
```

Rust Serde

<https://serde.rs/> - is a framework for serializing and deserializing Rust data structures efficiently and generically.

The Serde ecosystem consists of data structures that know how to serialize and deserialize themselves along with data formats that know how to serialize and deserialize other things. Serde provides the layer by which these two groups interact with each other, allowing any supported data structure to be serialized and deserialized using any supported data format.

JSON, Postcard, CBOR, YAML, MessagePack, TOML, Pickle, ROS, BSON, Avro, JSON5, URL, Envy, S-expressions, D-Bus, FlexBuffers, Bencode, DynamoDb, Hjson, ...

Roundtrip example

```
use serde::{Serialize, Deserialize};
```

```
#[derive(Serialize, Deserialize, Debug)]
```

```
struct Point {
```

```
    x: i32,
```

```
    y: i32,
```

```
}
```

```
fn main() {
```

```
    let point = Point { x: 1, y: 2 };
```

```
    // Convert the Point to a JSON string.
```

```
    let serialized = serde_json::to_string(&point).unwrap();
```

```
    // Prints serialized = {"x":1,"y":2}
```

```
    println!("serialized = {}", serialized);
```

```
    // Convert the JSON string back to a Point.
```

```
    let deserialized: Point = serde_json::from_str(&serialized).unwrap();
```

```
    // Prints deserialized = Point { x: 1, y: 2 }
```

```
    println!("deserialized = {:?}", deserialized);
```

```
}
```


qlog-dancer

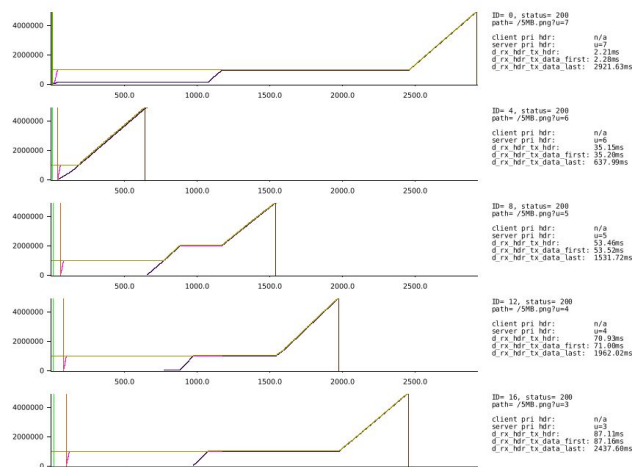
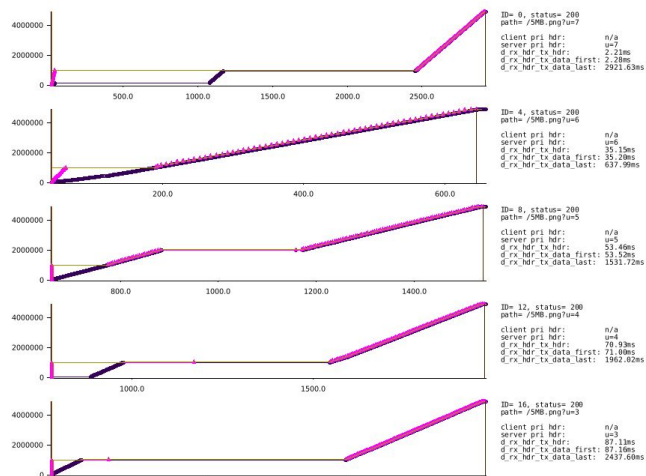
An internal tool. qlog crate powers deserialization. Application responsible for file handling, data munging and plotting (using [Plotters](#) library).

Builds on the shoulders of giants for congestion control oriented plots.




Qlog-dancer streams

Lots of spooky behavioural oddities happen when we cross the streams.



Qlog-dancer streams

Lots of spooky behavioural oddities happen when we cross the streams.



A real-world failure (1)

Reporter: Downloads are in Chrome are **behaving weirdly**. Seems it is worse with larger files.

Responder: **Logs or it didn't happen**. Also, got a repro?

Reporter: Downloads aren't captured in HAR files*. I can't get any other logs. I can describe the repro.

Responder: What about Chrome netlog?

Reporter: **What is Chrome netlog?**

Responder: `chrome://net-export/`

Reporter: OK, I captured you a netlog at the point the problem happened. While I was doing a dozen other things in the same browsing session.

Responder: D'oh. Don't worry, I'll recreate the repro and make my own logs.

* Narrator: we will never find out why HAR doesn't cover downloads. But also, for any HTTP/2 or HTTP/3 related error, HARs tend to be next to a bit rubbish.

A real-world failure (2)

```
#!/bin/bash
```

```
RAND_VAL=$RANDOM  
PROFILE_PATH="${HOME}/.temp-chrome-${RAND_VAL}"  
#PROFILE_PATH=$HOME/temp-chrome  
mkdir $PROFILE_PATH  
echo "using temporary profile at ${PROFILE_PATH}"
```

```
NETLOG_FILE="${HOME}/netlog_${RAND_VAL}.json"  
echo "logging netlog to ${NETLOG_FILE}"
```


```
# launch a fresh Chrome profile without annoying first-time checks and have it  
immediately start netlogging  
google-chrome --user-data-dir=$PROFILE_PATH --disable-fre --no-default-browser-check  
--no-first-run --log-net-log=$NETLOG_FILE --auto-open-devtools-for-tabs $1
```

```
# tidy up leftovers we don't want to persist
```

```
rm -r $PROFILE_PATH
```

A real-world failure (3)

- [Netlog viewer](#) crashes opening this jumbo log
- Manually data munge
- In file, lookup `HTTP2_SESSION_RECV_DATA` in the constants block and note it's number (it was 209 for me but it might vary?),
- do the same for `HTTP2_SESSION_RECV_RST_STREAM` (199 for me).
- Then delete the constants block from the netlog
- Run this query `"jq '.events[] | select(.params.stream_id==7 and .type==209) | [.time, .params.size] | @tsv' chrome-net-export-log-lucas.json > stream_data.tsv"`
- Plot the data in some tool like google sheets
- Find a way to never have to do this manually again.



Apply learnings from qlog to netlog

Qlog and netlog both JSON-ish formats

Qlog-dancer can parsing large files fast.
Reuse its framework.

Just need to see the netlog schema and
write some serde-compatible structures.



Reverse-engineering netlog to serde

```
#[derive(Serialize, Deserialize, Debug, Default)]
pub struct Http3DataFrameReceivedParams {
    pub payload_length: u64,
    pub stream_id: u64,
}

#[derive(Serialize, Deserialize, Debug, Default)]
pub struct Http3DataFrameReceivedEvent {
    pub params: Http3DataFrameReceivedParams,
}
```


Reverse-engineering netlog to serde

```
fn parse_netlog_h3_event(
    session: &mut Vec<(f32, netlog::Event)>,
    event_hdr: &netlog::EventHeader,
    event: &[u8], verbose: bool,
) {
    match event_hdr.ty_string.as_str() {
        "HTTP3_DATA_FRAME_RECEIVED" => {
            let ev: Http3DataFrameReceivedEvent =
                serde_json::from_slice(event).unwrap();
            session.push((
                netlog_time_delta(event_hdr),
                netlog::Event::H3(h3::Event::Http3DataFrameReceived(ev)),
            ));
        },
        _ =>()
    }
}
```

Summary

It would be great if there was one true logging format. But that's not realistic.

Well defined logging format can encourage an ecosystem of analysis tools. E.g. what is an endpoint producing and why is it doing that?

Tools compliment each other. Tools provide more value when they address user needs. Keeping tools internal reduces user base.

Is there any interest in some of the work presented? E.g., open sourcing of qlog-dancer or netlog library, trying to define schema for netlog or HAR?